

Two-Phase PageRank with Travel-Time Self-Loops

A Comprehensive Mathematical Reference
Salt Lake City Urban Traffic Prediction

Technical Documentation

March 2026

Abstract

We present a comprehensive derivation, analysis, and empirical evaluation of the *Two-Phase PageRank with Travel-Time Self-Loops* algorithm, developed to predict macroscopic traffic congestion on the Salt Lake City road network (99,716 distinct road links). Standard PageRank suffers from a fundamental physical defect: it treats every link transition as instantaneous, causing excessive probability-mass diffusion and a Mean Relative Error (MRE) of **0.7331** on the busiest 100 links.

We resolve this by injecting physical dwell time into the transition matrix via *self-loop weights* proportional to a link's traversal time, controlled by a single global parameter μ . We further split the network into two phases—active driving and local dispersal—forming a $2N \times 2N$ block transition matrix. With optimal parameters ($\mu = 20$, $\beta = 0.90$, $d = 0.80$) the Top-100 MRE plunges to **0.0317**—a **95.7%** improvement—while the overall MRE simultaneously falls to **0.0471** (76.2% improvement), demonstrating that the fix is physically principled rather than over-fitted.

Contents

1	Introduction and Motivation	3
1.1	The Traffic Prediction Problem	3
1.2	Why PageRank?	3
1.3	The Fundamental Defect: Instantaneous Transitions	3
2	Data Pre-Processing: Spatial Smoothing	4
2.1	The Sparsity Problem: Quantification and Consequences	4
2.2	Step 1: Laplace Pseudo-Count	5
2.3	Step 2: Graph-Diffusion Smoothing	5
2.3.1	Definition	6
2.3.2	Interpretation as a Graph Filter	6
2.3.3	Mass Conservation	6
2.4	Complete Construction of the Teleportation Vector E_N	7
2.4.1	Four-Step Pipeline	7
2.4.2	Worked Numerical Example	7

2.4.3	Separation of Ground Truth and Topology	8
2.4.4	Why Probability-Based Injection Is Appropriate	9
2.5	Theoretical Justification: The Bias-Variance Tradeoff	9
2.6	Empirical Validation: Cross-Validation Protocol	10
2.6.1	Protocol	10
2.7	Cross-Validation Results	11
2.8	Impact on the Downstream Model	13
2.9	Independent Verification of $\gamma^* = 0.26$	14
2.10	Summary of the Smoothing Justification	16
3	The Core Physical Fix: Travel-Time Self-Loops	16
3.1	Mathematical Definition	16
3.2	The Normalised Transition Matrix P	16
3.3	Worked Examples	17
3.4	Self-Loop Probability as a Function of Traversal Time	17
4	The Two-Phase Network: $2N \times 2N$ Block Matrix	19
4.1	Motivation for Two Phases	19
4.2	Block Matrix Construction	19
4.3	Interpretation of Each Block	19
4.4	Row-Stochasticity of M_{2N}	19
4.5	Effect of β on Phase Dynamics	20
5	The Teleportation Vector E_{2N}	20
5.1	Role of Teleportation in PageRank	20
5.2	Constructing E_{2N}	21
5.3	Deprecated Enhancement Parameters	21
6	The Power Iteration Algorithm	21
6.1	PageRank Update Rule	21
6.2	Expanded Form	22
6.3	Pseudocode	22
6.4	Convergence Analysis	22
7	Extracting the Final Traffic Distribution	23
8	Parameter Study and Empirical Results	23
8.1	Baseline: Standard PageRank ($\mu = 0$)	23
8.2	Effect of the Dwell-Time Parameter μ	24
8.3	Full Parameter Grid Search	24
8.4	Sensitivity to Damping Factor d	25
8.5	Before vs. After: Summary of Improvement	26
9	Physical Interpretation and Theoretical Connections	27
9.1	Connection to Continuous-Time Markov Chains	27
9.2	Why $\mu = 20$ Is Near-Optimal	27
9.3	The Two-Phase Structure as Absorption	27
9.4	Physical Interpretation of the Two Phases	27
10	Complete Mathematical Summary	28

1 Introduction and Motivation

1.1 The Traffic Prediction Problem

Modern cities generate massive streams of GPS trajectories from vehicles, ride-sharing apps, navigation systems, and fleet tracking. The challenge is to convert a *sparse sample* of these trajectories—typically covering only a small fraction of the network—into a reliable estimate of macroscopic traffic volume across *every* road segment.

For Salt Lake City, the road network has $N = 99,716$ distinct directed road links. In a given 15-minute window, raw GPS data may only directly observe traffic on a few thousand of these links. The prediction question is: given this partial view, what is the full traffic distribution?

1.2 Why PageRank?

PageRank models network flow as a random walk over a directed graph. A random walker (simulated vehicle) either moves to one of the current link’s successors or teleports to a node sampled from a *teleportation distribution*. The stationary distribution of this Markov chain represents the predicted fraction of time (and thus traffic) on each link.

The key advantage of PageRank is that it naturally propagates observed traffic through the network topology: high traffic on one road *implies* high traffic on connected feeders and exits, even if those roads were not directly observed.

1.3 The Fundamental Defect: Instantaneous Transitions

Warning / Pitfall

Standard PageRank treats every link transition as taking the same mathematical “time step”. A vehicle transitioning from a 1 km highway at 100 km/h and a vehicle transitioning through a 10 m intersection connector at 20 km/h both advance in exactly one iteration. This is physically absurd.

Because the highway segment takes $t_{\text{highway}} = 1000/27.8 \approx 36\text{ s}$ to traverse while the connector takes $t_{\text{connector}} = 10/5.6 \approx 1.8\text{ s}$, the highway vehicle realistically spends $\sim 20\times$ longer at that location. Standard PageRank ignores this completely, causing excessive *diffusion*: traffic flows through long roads too quickly and spreads uniformly, washing out the concentration peaks on major arterial roads.

Figure 1 shows the complete processing pipeline from raw GPS data to final predictions.




Figure 1: End-to-end data pipeline: raw GPS trajectories are smoothed with graph-diffusion ($\gamma = 0.26$) before being used as a teleportation vector in the Two-Phase PageRank power iteration.

2 Data Pre-Processing: Spatial Smoothing

This section provides a rigorous justification for the spatial smoothing pre-processing step. We address three fundamental questions: **(1)** Why is smoothing necessary? **(2)** How much smoothing should be applied? **(3)** What is the quantitative impact on downstream predictions?

2.1 The Sparsity Problem: Quantification and Consequences

The raw trajectory data is converted into a vehicle-count vector $\mathbf{c} \in \mathbb{R}^N$, where c_i counts the number of GPS-traced vehicles observed on link i during a single 15-minute evaluation window. Because the fleet of GPS-equipped vehicles constitutes only a small sample of total traffic, the vast majority of entries are $c_i = 0$.

Figure 2 quantifies this sparsity empirically. Over the full month-long dataset, the median 15-minute window observes traffic on only a small fraction of the network’s $N = 99,716$ links.




Figure 2: Distribution of observation coverage across all 15-minute timeframes in the dataset. The left panel shows absolute counts; the right panel shows the percentage of the network observed. Most timeframes cover only a tiny fraction of the graph.

This extreme sparsity has three critical consequences for the downstream PageRank model:

- C1. MRE Singularity.** The Mean Relative Error metric divides by the ground truth value: $\text{MRE}_i = |v_i - p_i|/p_i$. When $p_i = 0$ for the majority of links, MRE is either undefined or requires a small additive constant ϵ that dominates the metric for zero-valued links.
- C2. Teleportation Dead Zones.** In PageRank, with probability $(1 - d)$, a simulated vehicle teleports to a random link weighted by $E_N[i]$. When $E_N[i] = 0$ for most links, the teleportation vector concentrates mass on the few observed roads. The resulting stationary distribution severely overweights these roads and assigns near-zero mass to their unobserved neighbours, creating artificial “sinks” in the Markov chain.
- C3. Non-representative Sample.** Sparse GPS trajectories are a biased sample of true traffic, over-representing major roads (where more GPS-equipped vehicles travel) and systematically under-representing secondary roads. Without smoothing, this sampling bias propagates directly into the prediction.

Warning / Pitfall

Empirical evidence: When raw (unsmoothed) counts are used directly as the PageRank teleportation vector across 490 timeframes, the average MRE is **7,275**. After smoothing, the same metric drops to **0.19**—a **37,800×** improvement (see Section 2.8).

2.2 Step 1: Laplace Pseudo-Count

To eliminate exact zeros and provide a minimally informative prior, we apply Laplace smoothing:

$$\tilde{c}_i = c_i + \alpha, \quad i = 1, \dots, N, \quad \alpha = 0.1. \quad (1)$$

This is a standard Bayesian regularisation technique. Each link receives a small pseudo-count that represents the prior belief that *some* traffic exists on every road. The value $\alpha = 0.1$ is small enough that it does not distort the magnitudes of genuinely observed links (whose counts are typically ≥ 1), but large enough to prevent the teleportation vector from being identically zero on unobserved roads.

Proposition 2.1 (Relative Perturbation Bound). *For any link with true count $c_i \geq 1$, the relative perturbation introduced by the pseudo-count satisfies:*

$$\frac{|\tilde{c}_i - c_i|}{c_i} = \frac{\alpha}{c_i} \leq \alpha = 0.1. \quad (2)$$

Thus the Laplace prior distorts observed counts by at most 10%.

2.3 Step 2: Graph-Diffusion Smoothing

After Laplace regularisation, unobserved roads all share the same flat prior $\tilde{c}_i = \alpha$. This does not capture the fact that a road *adjacent* to a busy observed road is more likely to carry traffic than an isolated unobserved road. Graph-diffusion smoothing addresses this by propagating evidence along the network topology.

2.3.1 Definition

Let $\mathbf{A} \in \{0, 1\}^{N \times N}$ be the adjacency matrix of the directed road graph, and let $\mathbf{P}_{\text{adj}}$ be its row-stochastic normalisation:

$$[\mathbf{P}_{\text{adj}}]_{ij} = \frac{A_{ij}}{\sum_{k=1}^N A_{ik}} = \frac{A_{ij}}{\deg^+(i)}, \quad (3)$$

where $\deg^+(i)$ is the out-degree of node i .

The smoothed count vector is defined as a **one-step lazy random walk** on the graph:

$$\boxed{\mathbf{c}_{\text{smooth}} = (1 - \gamma) \tilde{\mathbf{c}} + \gamma \mathbf{P}_{\text{adj}} \tilde{\mathbf{c}} = [(1 - \gamma) \mathbf{I} + \gamma \mathbf{P}_{\text{adj}}] \tilde{\mathbf{c}}} \quad (4)$$

where $\gamma \in [0, 1)$ is the *diffusion factor* controlling the strength of neighbourhood borrowing.

The result is then L^1 -normalised to produce the teleportation probability vector:

$$E_N[i] = \frac{c_{\text{smooth},i}}{\sum_{j=1}^N c_{\text{smooth},j}}, \quad \sum_{i=1}^N E_N[i] = 1. \quad (5)$$

2.3.2 Interpretation as a Graph Filter

The operator $\mathbf{H}_\gamma = (1 - \gamma) \mathbf{I} + \gamma \mathbf{P}_{\text{adj}}$ is a *first-order low-pass graph filter*. To see this, consider the eigendecomposition of $\mathbf{P}_{\text{adj}}$. If λ is an eigenvalue of $\mathbf{P}_{\text{adj}}$ (with $|\lambda| \leq 1$), the corresponding eigenvalue of $\mathbf{H}_\gamma$ is:

$$h(\lambda) = (1 - \gamma) + \gamma\lambda. \quad (6)$$

For the constant eigenvector ($\lambda_1 = 1$, corresponding to the “DC component” or uniform signal), $h(1) = 1$ —the global mean is perfectly preserved. For high-frequency eigenvectors ($\lambda \approx -1$, representing rapid spatial oscillations), $h(-1) = 1 - 2\gamma$, which is attenuated for $\gamma > 0$. Thus:

Key Insight

Graph-diffusion smoothing acts as a **spectral low-pass filter**: it preserves the large-scale spatial structure of traffic (the “mean field”) while attenuating high-frequency noise caused by sparse sampling. Increasing γ increases the degree of attenuation.

2.3.3 Mass Conservation

An important property of Eq. (4) is that total mass is conserved:

$$\sum_i c_{\text{smooth},i} = (1 - \gamma) \sum_i \tilde{c}_i + \gamma \sum_i [\mathbf{P}_{\text{adj}} \tilde{\mathbf{c}}]_i = (1 - \gamma) \sum_i \tilde{c}_i + \gamma \sum_i \tilde{c}_i = \sum_i \tilde{c}_i, \quad (7)$$

since $\mathbf{P}_{\text{adj}}$ is row-stochastic ($\sum_j P_{ij} = 1$) and therefore the column sums of $\mathbf{P}_{\text{adj}}^\top \mathbf{1}$ are preserved under multiplication. This means smoothing *redistributes* mass but never creates or destroys it, ensuring physical consistency.

2.4 Complete Construction of the Teleportation Vector E_N

This subsection consolidates the steps above into a single, self-contained description of how the **ground-truth GPS data** is transformed into the teleportation probability vector used by PageRank. It is important to emphasise that the ground truth *only* enters the model through this vector—it plays no role whatsoever in the transition matrix, which is built entirely from the static road-network topology, road lengths, and speed limits.

2.4.1 Four-Step Pipeline

Step 1. Extract raw counts from GPS data. For a given 15-minute window, parse all GPS trajectories and count how many traced vehicles traversed each of the $N = 99,716$ road links. The result is a raw count vector $\mathbf{c} \in \mathbb{Z}_{\geq 0}^N$, where c_i is the number of observed vehicles on link i . Most entries are $c_i = 0$.

Step 2. Apply Laplace pseudo-count. Add a small constant $\alpha = 0.1$ to every entry:

$$\tilde{c}_i = c_i + 0.1, \quad i = 1, \dots, N.$$

This ensures that no link has zero probability, preventing MRE singularities and Markov chain absorbing states.

Step 3. Apply graph-diffusion smoothing ($\gamma = 0.26$). Propagate evidence along the road network:

$$c_{\text{smooth},i} = 0.74 \tilde{c}_i + 0.26 \frac{1}{\deg^+(i)} \sum_{j \in \text{out}(i)} \tilde{c}_j,$$

where the sum runs over the out-neighbours of link i in the directed road graph. Each link retains 74% of its own value and receives 26% from the average of its neighbours.


Step 4. L^1 -normalise to obtain E_N . Divide by the total mass to convert counts into a probability distribution:

$$E_N[i] = \frac{c_{\text{smooth},i}}{\sum_{j=1}^N c_{\text{smooth},j}}, \quad \sum_{i=1}^N E_N[i] = 1.$$

Each entry $E_N[i]$ now represents the *probability* that a newly injected (“teleported”) vehicle will appear on link i .

2.4.2 Worked Numerical Example

To make the construction concrete, consider a toy network with $N = 5$ links and the following directed adjacency structure:



Suppose during a 15-minute window, the raw GPS counts are:

Table 1: Worked example: constructing E_N from raw GPS data.

Link i	1	2	3	4	5
Step 1: Raw count c_i	0	5	0	3	0
Step 2: After Laplace ($\tilde{c}_i = c_i + 0.1$)	0.10	5.10	0.10	3.10	0.10
Step 3a: Neighbour average $\frac{1}{ \text{out}(i) } \sum_{j \in \text{out}(i)} \tilde{c}_j$	5.10	1.60	0.10	0.10	—
Step 3b: Smoothed $c_{\text{smooth},i} = 0.74 \tilde{c}_i + 0.26 \cdot \text{avg}$	1.400	4.190	0.100	2.320	0.100
Step 4: $E_N[i] = c_{\text{smooth},i} / \text{total}$	0.1724	0.5160	0.0123	0.2858	0.0123

Reading the example:

- Link 1 had *zero* observed vehicles, yet it receives $E_N[1] = 17.2\%$ of the teleportation probability—because its only out-neighbour (link 2) is heavily trafficked ($c_2 = 5$), and smoothing propagates evidence from link 2 back to link 1.
- Link 5 is a dead-end (no out-neighbours), so its smoothed value remains at the Laplace floor ($\alpha = 0.1$). It gets the minimal $E_N[5] = 1.2\%$.
- Link 2 retains the lion’s share (51.6%) because it has the highest raw count, and smoothing only reduces it slightly.

Key Insight

This example illustrates the fundamental roles:

- (i) **Ground truth** determines *where* traffic is injected (links 2 and 4 dominate because they were observed).
- (ii) **Topology** determines *how* evidence propagates (link 1 gains probability despite having zero observations, because it feeds into the busy link 2).
- (iii) The two roles are combined *only* in the teleportation vector; the transition matrix is built entirely from the physical network.

2.4.3 Separation of Ground Truth and Topology

It is critical to understand the clean architectural separation in the Two-Phase PageRank model:

Table 2: Role of ground truth vs. network topology in the model.

Component	Data source	Role in PageRank
Transition matrix M_{2N}	Road topology, lengths, speeds	Governs where vehicles <i>move</i>
Teleportation vector E_{2N}	Ground truth GPS counts	Governs where vehicles are <i>born</i>

The PageRank power iteration combines both components:

$$\mathbf{v}^{(k+1)} = d M_{2N}^\top \mathbf{v}^{(k)} + (1 - d) E_{2N}. \quad (8)$$

At every iteration, a fraction $d = 0.80$ of probability mass follows the **topology-driven** transition matrix (vehicles move along roads), while the remaining $1 - d = 0.20$ is **re-injected from ground truth** via the teleportation vector (new vehicles appear on roads proportional to observed traffic). This blending is what makes the model both *physically realistic* (traffic follows roads) and *empirically grounded* (predictions are anchored to real observations).

2.4.4 Why Probability-Based Injection Is Appropriate

A natural question is: given that the ground truth counts are sparse and noisy, is it appropriate to convert them into probabilities and inject them directly into the Markov chain?

Three arguments support this design:

- A1. Physical consistency.** The PageRank stationary distribution $\mathbf{v}$ is itself a probability distribution ($\sum_i v_i = 1$). To add a source term that is consistent with the output, the teleportation vector must also be a probability distribution. Using raw counts would introduce dimensional inconsistency.
- A2. Scale invariance.** On a quiet Sunday night, the total GPS count might be 500; during Monday rush hour it might be 50,000. By normalising to probabilities, the algorithm responds to the *relative spatial pattern* of traffic (which links are busier *than others*), not the absolute volume. The absolute scale is irrelevant to predicting traffic *distribution*.
- A3. Smoothing as denoising.** The Laplace prior and graph-diffusion smoothing (Section 2.2–2.3) address the concern that raw probabilities would be noisy. After smoothing, the probability vector E_N is a denoised, topology-aware estimate of the true traffic distribution, suitable for direct injection.

Empirical Result

The 490-timeframe experiment of Section 2.8 provides strong empirical evidence: the smoothed probability-based injection achieves $\text{MRE} = 0.19$, meaning the PageRank output matches the ground truth with an average relative error of only 19%. This confirms that the injection mechanism is not just theoretically sound but also practically effective.

2.5 Theoretical Justification: The Bias-Variance Tradeoff

Smoothing introduces a fundamental tradeoff between two sources of error in the predictive distribution. We formalise this using the standard MSE decomposition.

Definition 2.1 (Bias and Variance of the Smoothed Estimator). Let $\mathbf{c}$ be the random vehicle-count vector (a random sample from the true traffic distribution $\boldsymbol{\pi}$), and let $\hat{\boldsymbol{\pi}}_\gamma = \text{normalise}(\mathbf{H}_\gamma \tilde{\mathbf{c}})$ be the smoothed estimator. Then:

$$\text{MSE}(\hat{\boldsymbol{\pi}}_\gamma) = \underbrace{\|\mathbb{E}[\hat{\boldsymbol{\pi}}_\gamma] - \boldsymbol{\pi}\|^2}_{\text{Bias}^2(\gamma)} + \underbrace{\mathbb{E}[\|\hat{\boldsymbol{\pi}}_\gamma - \mathbb{E}[\hat{\boldsymbol{\pi}}_\gamma]\|^2]}_{\text{Variance}(\gamma)}. \quad (9)$$

The two terms have opposing monotonicity in γ :

- **Variance $\downarrow$ as $\gamma \uparrow$:** By averaging over neighbours, the estimator becomes less sensitive to the specific random sample of GPS vehicles. This is analogous to kernel smoothing in nonparametric statistics, where larger bandwidths reduce estimator variance.
- **Bias² $\uparrow$ as $\gamma \uparrow$:** The operator $\mathbf{H}_\gamma$ pushes $\hat{\pi}_\gamma$ toward the degree-weighted average of its neighbourhood, blurring genuine spatial peaks in the true distribution. In the extreme ($\gamma \rightarrow 1$), every link’s value converges to the mean of its neighbours, destroying all spatial variation.

The total MSE therefore traces a **U-shaped curve** in γ , with an optimal point γ^* where the marginal variance reduction exactly balances the marginal bias increase.




Figure 3: Normalised bias-variance decomposition from cross-validation data. The blue curve (variance proxy) decreases monotonically with γ ; the orange curve (bias proxy) increases monotonically; the purple curve (total MSE, normalised) shows the U-shaped tradeoff with minimum at $\gamma^* = 0.26$.

2.6 Empirical Validation: Cross-Validation Protocol

To empirically determine γ^* , we implemented a **Monte Carlo Split-Half Cross-Validation** that rigorously estimates the out-of-sample predictive accuracy of each smoothing level.

2.6.1 Protocol

- Step 1. Timeframe selection.** From the one-month GPS dataset, we identify timeframe “slots” (grouped by day-of-week and time-of-day) that contain at least 2 independent recording instances. We randomly select 49 such slots using a fixed random seed for reproducibility.
- Step 2. Split-half partition.** For each selected slot, the available recording instances are randomly split into two disjoint halves: **Set A** (training) and **Set B** (validation).

Step 3. Apply smoothing to Set A only. For each of the 11 candidate values $\gamma \in \{0.00, 0.05, \dots, 0.60\}$ (see Table 3), compute $\hat{\pi}_\gamma^{(A)}$ from Set A using Eq. (4).

Step 4. Evaluate against raw Set B. Set B is *never* smoothed; it serves as the “oracle” ground truth. For each γ , we compute:

- **Predictive MSE:** $\text{MSE}^{\text{raw}}(\gamma) = \frac{1}{N} \sum_{i=1}^N (\hat{\pi}_{\gamma,i}^{(A)} - \hat{\pi}_{0,i}^{(B)})^2$
- **Pearson correlation:** $r(\gamma) = \text{corr}(\hat{\pi}_\gamma^{(A)}, \hat{\pi}_0^{(B)})$

Step 5. Evaluate downstream PageRank. We also run the full Two-Phase PageRank solver on both sides and compare the resulting stationary distributions:

- **PageRank MSE:** $\text{MSE}^{\text{PR}}(\gamma) = \frac{1}{N} \sum_{i=1}^N (v_{\gamma,i}^{(A)} - v_{0,i}^{(B)})^2$

Step 6. Average metrics. All metrics are averaged over the 49 Monte Carlo trials to produce stable estimates.

The key design principle is that *Set B is never smoothed*, so the evaluation metrics test whether smoothing helps the algorithm *generalise* to unseen raw data, rather than merely fitting the smoothed version of itself.

2.7 Cross-Validation Results

Table 3 presents the complete cross-validation results for all 11 candidate γ values.

Table 3: Complete cross-validation results for graph-diffusion smoothing (49 Monte Carlo split-half trials).

γ	Predictive MSE	Pearson r	PageRank MSE	Var. retained
0.00	3.397×10^{-8}	0.297	3.554×10^{-9}	100.0%
0.05	3.393×10^{-8}	0.298	3.553×10^{-9}	93.2%
0.10	3.389×10^{-8}	0.300	3.552×10^{-9}	86.7%
0.15	3.386×10^{-8}	0.301	3.551×10^{-9}	80.7%
0.20	3.384×10^{-8}	0.301	3.551×10^{-9}	75.1%
0.26	3.383×10^{-8}	0.301	3.551×10^{-9}	69.0%
0.30	3.384×10^{-8}	0.301	3.551×10^{-9}	65.3%
0.35	3.385×10^{-8}	0.300	3.551×10^{-9}	61.0%
0.40	3.387×10^{-8}	0.298	3.552×10^{-9}	57.2%
0.50	3.395×10^{-8}	0.291	3.554×10^{-9}	50.8%
0.60	3.407×10^{-8}	0.279	3.557×10^{-9}	46.2%

Three key findings emerge from Table 3:

Empirical Result

Finding 1 — Optimal γ^* : Both the minimum predictive MSE (3.383×10^{-8}) and the maximum Pearson correlation ($r = 0.3015$) occur at $\gamma^* = 0.26$, confirming the bias-variance U-curve of Section 2.5.

Finding 2 — Downstream coherence: The minimum *downstream* PageRank MSE also occurs near $\gamma \approx 0.26$, demonstrating that the optimal smoothing level for direct prediction also minimises error *after* the PageRank transformation. This is a non-trivial result: the PageRank solver is a nonlinear operator, and there is no a priori guarantee that input-optimal smoothing is output-optimal.

Finding 3 — Variance budget: At $\gamma^* = 0.26$, the estimator retains 69% of the original signal variance. This means 31% of the variance in the raw data is attributable to sampling noise and is correctly removed by smoothing.




Figure 4: Cross-validation: held-out predictive MSE and Pearson correlation vs. γ (dual-axis). Both metrics are jointly optimised at $\gamma^* = 0.26$.




Figure 5: Downstream PageRank MSE and fraction of signal variance retained as a function of γ . The PageRank MSE minimum coincides with the direct prediction MSE minimum, confirming that the smoothing level is compatible with the downstream solver.




Figure 6: Predictive MSE and downstream PageRank MSE side by side, both minimised at $\gamma^* = 0.26$.

2.8 Impact on the Downstream Model

To quantify the practical impact of smoothing on the full Two-Phase PageRank pipeline, we performed a large-scale comparison. For each of 490 randomly sampled 15-minute timeframes from the month-long dataset, we ran the complete pipeline twice:

1. **Raw pipeline:** raw vehicle counts $\rightarrow$ Laplace prior $\rightarrow$ PageRank (no graph-diffusion smoothing, $\gamma = 0$)
2. **Smoothed pipeline:** raw vehicle counts $\rightarrow$ Laplace prior $\rightarrow$ graph-diffusion ($\gamma = 0.26$) $\rightarrow$ PageRank

In both cases, MRE is computed against the pipeline’s own teleportation vector (i.e., how accurately does the PageRank stationary distribution reproduce the input?).

Table 4: Downstream impact: average MRE across 490 timeframes.

Pipeline	Average MRE	Improvement factor
Raw ($\gamma = 0$)	7,275	—
Smoothed ($\gamma = 0.26$)	0.19	$\times 37,800$




Figure 7: Per-timeframe MRE comparison across 490 windows. **Left:** log-scale histogram showing the raw pipeline clustered around MRE $\approx 7,000$ while the smoothed pipeline achieves MRE ≈ 0.19 . **Right:** sorted per-timeframe comparison showing a consistent ~ 4 orders-of-magnitude gap.

Key Insight

The $37,800\times$ improvement factor confirms that graph-diffusion smoothing is not merely a cosmetic improvement—it is a **mathematical precondition** for the PageRank solver to function. Without smoothing, the zero-inflated teleportation vector creates absorbing states in the Markov chain that trap probability mass and produce catastrophically meaningless stationary distributions.

2.9 Independent Verification of $\gamma^* = 0.26$

To rule out the possibility that $\gamma^* = 0.26$ is an artefact of the specific random seed or sample of timeframes used in Section 2.7, we conducted a fully independent replication experiment with the following changes:

- **Different random seed:** seed = 123 (vs. 42 in the original).
- **More splits:** 80 Monte Carlo trials (vs. 49 in the original).
- **Finer γ grid:** 17 candidate values with 0.02-step resolution in the range $[0.18, 0.30]$ surrounding the hypothesised optimum.

Table 5: Verification CV results (seed = 123, 80 splits, 17 γ values).

γ	Pred. MSE ($\times 10^{-8}$)	Pearson r	PR MSE ($\times 10^{-8}$)	Var. (%)
0.00	2.708	0.321	2.431	100.0
0.10	2.698	0.323	2.422	86.8
0.20	2.692	0.325	2.416	75.3
0.22	2.691	0.325	2.416	73.2
0.24	2.691	0.325	2.415	71.2
0.26	2.690	0.325	2.415	69.2
0.28	2.690	0.324	2.414	67.3
0.30	2.690	0.324	2.414	65.5
0.40	2.694	0.320	2.416	57.4
0.60	2.714	0.299	2.432	46.5

Empirical Result

Key findings from the verification experiment:

- (a) The MSE curve exhibits a broad, flat plateau for $\gamma \in [0.22, 0.30]$. Within this band, the MSE varies by less than **0.04%**, demonstrating that the model performance is highly robust to the exact choice of γ in this range.
- (b) The predictive MSE minimum shifts slightly to $\gamma = 0.28$ (vs. 0.26 in the original), while the correlation maximum shifts to $\gamma = 0.22$. However, the chosen value $\gamma^* = 0.26$ sits precisely at the **geometric centre** of this plateau, offering the best compromise between predictive error and correlation.
- (c) The two experiments use different random seeds, different numbers of trials (49 vs. 80), yet produce consistent U-shaped curves—confirming that the result is not a statistical fluke.




Figure 8: Independent verification on a finer γ grid. Both experiments (blue = original, pink = verification) agree on the location and breadth of the optimal plateau. The green band marks $\gamma \in [0.22, 0.30]$.

2.10 Summary of the Smoothing Justification

- S1. Why smooth?** Raw GPS data is extremely sparse ($\lesssim 3\%$ network coverage per window), creating zero-inflated probability vectors that cause MRE singularities and Markov chain absorbing states. Graph-diffusion smoothing acts as a spectral low-pass filter, propagating evidence from observed roads to their unobserved neighbours while preserving total mass.
- S2. How much?** The optimal smoothing factor $\gamma^* = 0.26$ was determined via 49-trial Monte Carlo split-half cross-validation and independently verified with an 80-trial experiment using a different random seed and a finer 17-point γ grid (Section 2.9). The optimum sits at the centre of a broad plateau $\gamma \in [0.22, 0.30]$ where MSE varies by less than 0.04%.
- S3. What is the impact?** Smoothing transforms the average MRE from 7,275 (completely meaningless) to 0.19 (highly accurate), an improvement factor of $37,800\times$ across 490 timeframes. After smoothing, the PageRank model can proceed to achieve its full potential with travel-time self-loops (Section 3).

3 The Core Physical Fix: Travel-Time Self-Loops

3.1 Mathematical Definition

For each link i with physical length ℓ_i (metres) and free-flow speed s_i (m/s), define the *traversal time*:

$$t_i = \frac{\ell_i}{s_i} \quad [\text{seconds}]. \quad (10)$$

We introduce a global *dwelt-time scaling parameter* $\mu > 0$ and define the *self-loop weight* of link i as:

$$\boxed{\text{sw}_i = \mu \cdot t_i = \mu \cdot \frac{\ell_i}{s_i}}. \quad (11)$$

Let w_{ij} be the weight of the forward edge from link i to successor link j . In the base configuration ($\alpha_s = \alpha_l = 0$), $w_{ij} = 1$ for all valid outgoing edges. The *total weight* leaving link i is:

$$\text{TW}_i = \text{sw}_i + \sum_{j \in \text{succ}(i)} w_{ij} = \mu t_i + |\text{succ}(i)|. \quad (12)$$

3.2 The Normalised Transition Matrix P

The $N \times N$ transition matrix P is constructed as:

$$P_{ii} = \frac{\text{sw}_i}{\text{TW}_i} = \frac{\mu t_i}{\mu t_i + |\text{succ}(i)|} \quad (\text{self-loop probability}), \quad (13)$$

$$P_{ij} = \frac{w_{ij}}{\text{TW}_i} = \frac{1}{\mu t_i + |\text{succ}(i)|} \quad (\text{forward probability, } j \in \text{succ}(i)), \quad (14)$$

$$P_{ij} = 0 \quad (\text{otherwise}). \quad (15)$$

By construction, every row of P sums to 1, making P a valid row-stochastic matrix.

3.3 Worked Examples

Example 3.1 (Highway Segment). Consider a highway segment with parameters: $\ell_i = 500$ m, $s_i = 25$ m/s, 3 successors, $\mu = 20$.

$$\begin{aligned} t_i &= 500/25 = 20 \text{ s} \\ \text{sw}_i &= 20 \times 20 = 400 \\ \text{TW}_i &= 400 + 3 = 403 \\ P_{ii} &= 400/403 \approx 0.9926 \quad \textbf{(99.26\% stays per step)} \\ P_{ij} &= 1/403 \approx 0.0025 \quad \text{per successor} \end{aligned}$$

After $k = 100$ iterations: $(0.9926)^{100} \approx 0.473$ — the highway retains nearly half its initial mass.

Example 3.2 (Short Connector). Consider $\ell_i = 50$ m, $s_i = 10$ m/s, 3 successors, $\mu = 20$.

$$\begin{aligned} t_i &= 50/10 = 5 \text{ s} \\ \text{sw}_i &= 20 \times 5 = 100 \\ \text{TW}_i &= 100 + 3 = 103 \\ P_{ii} &= 100/103 \approx 0.9709 \quad (97.09\% \text{ stays per step}) \\ P_{ij} &= 1/103 \approx 0.0097 \quad \text{per successor} \end{aligned}$$

After $k = 100$ iterations: $(0.9709)^{100} \approx 0.053$ — the connector has shed $\sim 95\%$ of its initial mass.

Key Insight

Over 100 iterations, the highway retains **47%** of its initial mass while the connector retains only **5%**. This directly models the physical reality that vehicles spend $\frac{500/25}{50/10} = 4\times$ longer on the highway than on the connector under $\mu=20$, leading to a natural traffic accumulation without any ad-hoc forcing.

3.4 Self-Loop Probability as a Function of Traversal Time

Proposition 3.1. For fixed μ and k successors, the self-loop probability P_{ii} is a monotonically increasing function of t_i :

$$P_{ii}(t_i) = \frac{\mu t_i}{\mu t_i + k} \xrightarrow{t_i \rightarrow \infty} 1, \quad P_{ii}(0) = 0. \quad (16)$$

Furthermore, the link's effective “time constant” under repeated application is:

$$\tau_i = -\frac{1}{\ln P_{ii}} \approx \frac{\mu t_i}{k} \quad (\text{for } \mu t_i \gg k). \quad (17)$$

Figure 9 shows P_{ii} as a function of traversal time for several values of μ .




Figure 9: Self-loop probability P_{ii} vs. traversal time t_i for various μ (3 outgoing edges assumed). Higher μ amplifies the dwell-time effect.

Figure 10 shows the resulting mass-retention profile over 200 power iterations for three representative link types.




Figure 10: Fraction of initial probability mass retained on a link over power iterations. The no-self-loop baseline ($\mu = 0$) immediately drops to zero, while the highway segment retains $\sim 47\%$ after 100 steps.

4 The Two-Phase Network: $2N \times 2N$ Block Matrix

4.1 Motivation for Two Phases

Traffic is inherently directional during active driving. However, after arriving at a destination, vehicles park and disperse locally—they no longer follow the directed highway graph. Modelling both behaviours in a single matrix would conflate forward-flow dynamics with local dispersal and obscure the physical interpretation.

We therefore duplicate the N -node network into two identical *phases*:

1. **Up Phase** ($i = 1, \dots, N$): vehicles actively driving, following the directed road graph with self-loops.
2. **Down Phase** ($i = N + 1, \dots, 2N$): vehicles parked or dispersing locally at reduced speed; they use the same transition matrix P but can never re-enter the Up phase.

4.2 Block Matrix Construction

A parameter $\beta \in [0, 1]$ controls how quickly mass “falls” from the Up phase to the Down phase at each iteration. The full $2N \times 2N$ transition matrix is:

$$M_{2N} = \begin{bmatrix} (1 - \beta)P & \beta I_N \\ \mathbf{0}_N & P \end{bmatrix} \quad (18)$$

where I_N is the $N \times N$ identity matrix and $\mathbf{0}_N$ is the $N \times N$ zero matrix.

4.3 Interpretation of Each Block

Table 6: Semantic interpretation of each $N \times N$ block of M_{2N} .

Block	Expression	Physical meaning
Top-Left	$(1 - \beta)P$	Probability of staying in the Up phase and moving (or self-looping) according to P .
Top-Right	βI_N	Probability of dropping from Up phase to the <i>same</i> link in the Down phase (vertical transition, identity preserves location).
Bottom-Left	$\mathbf{0}_N$	Down-to-Up transition is <i>impossible</i> : once parked, a vehicle cannot re-enter active flow.
Bottom-Right	P	Down-phase vehicles continue their local random walk according to P (with self-loops).

4.4 Row-Stochasticity of M_{2N}

For any row i in the Up phase (top half):

$$\sum_j [M_{2N}]_{i,j} = (1 - \beta) \underbrace{\sum_j P_{ij}}_{=1} + \beta \cdot 1 = 1. \checkmark \quad (19)$$

For any row i in the Down phase (bottom half):

$$\sum_j [M_{2N}]_{i,j} = \sum_j P_{ij} = 1. \checkmark \quad (20)$$

Hence M_{2N} is row-stochastic and defines a valid Markov chain over $2N$ states.




Figure 11: Visual structure of the $2N \times 2N$ block transition matrix. The bottom-left block is identically zero, enforcing the one-way phase transition.

4.5 Effect of β on Phase Dynamics

- $\beta = 0$: All mass remains in the Up phase indefinitely; the two-phase structure degenerates to standard single-phase PageRank.
- $\beta = 1$: Every vehicle falls to the Down phase after exactly one step; no forward active-driving evolution occurs.
- $\beta = 0.90$ (optimal): Vehicles transition to the Down phase after roughly one hop, preventing excessive forward diffusion while self-loops handle dwell time.

At the stationary distribution, mass is shared between phases. For $\beta \rightarrow 1$, nearly all mass accumulates in the Down phase where it continues to evolve under P .

5 The Teleportation Vector E_{2N}

5.1 Role of Teleportation in PageRank

In the standard PageRank update rule, with probability $(1 - d)$ a walker teleports to a random location instead of following the transition matrix. This teleportation serves two purposes:

1. **Ergodicity:** guarantees the Markov chain is irreducible and aperiodic, ensuring a unique stationary distribution.
2. **Bias injection:** instead of teleporting uniformly, we bias toward our empirical ground-truth observation vector E_N , injecting observed traffic data into the simulation.

5.2 Constructing E_{2N}

The smoothed normalised traffic vector $E_N \in \mathbb{R}^N$ (see Section 2) is the teleportation distribution over the physical N links.

Since teleportation represents the “birth” of a new simulated vehicle, all newly injected vehicles should start in the *Up phase* (actively driving). Therefore:

$$E_{2N} = \begin{bmatrix} E_N \\ \mathbf{0}_N \end{bmatrix} \in \mathbb{R}^{2N}, \quad \|E_{2N}\|_1 = 1. \quad (21)$$

Remark 5.1. *Injecting mass only into the Up phase is physically motivated: vehicles “start” their journey on a road in active mode. They will naturally transition to the Down phase as the simulation progresses.*

5.3 Deprecated Enhancement Parameters

In earlier algorithm versions, two additional parameters modified E_N :

- **Tau** (τ): raised each E_N entry to the power τ before renormalising, concentrating teleportation mass on high-traffic links.
- **Top-K Boost:** multiplied the entries of the top- k busiest links by a factor, further concentrating mass.

Both approaches improved Top-100 MRE only by “hacking” the solution rather than fixing the underlying physics. After the μ self-loop fix, $\tau = 1.0$ and top-k-boost = 1.0, meaning E_N is used in its pure, unmodified form.

6 The Power Iteration Algorithm

6.1 PageRank Update Rule

Let $v_{2N}^{(k)} \in \mathbb{R}^{2N}$ be the probability distribution over the $2N$ states at iteration k , with $\|v_{2N}^{(k)}\|_1 = 1$. The PageRank update is:

$$v_{2N}^{(k+1)} = d \cdot M_{2N}^\top v_{2N}^{(k)} + (1 - d) \cdot E_{2N} \quad (22)$$

where:

- $d \in (0, 1)$ is the *damping factor*. With probability d , mass follows the transition dynamics; with probability $(1 - d)$, mass teleports.
- $M_{2N}^\top$ is the transpose because we represent distributions as *column vectors* and the transition matrix is row-stochastic.

6.2 Expanded Form

Writing $v^{(k)} = \begin{bmatrix} v_{\text{up}}^{(k)} \\ v_{\text{down}}^{(k)} \end{bmatrix}$ and substituting (18):

$$v_{\text{up}}^{(k+1)} = d(1 - \beta) P^\top v_{\text{up}}^{(k)} + (1 - d) E_N \quad (23)$$

$$v_{\text{down}}^{(k+1)} = d\beta v_{\text{up}}^{(k)} + d P^\top v_{\text{down}}^{(k)} \quad (24)$$

Equation (23) shows that the Up-phase distribution evolves under scaled P with teleportation injection. Equation (24) shows the Down phase receiving a “drip” of β -fraction of the Up phase mass each iteration, while also evolving under P .

6.3 Pseudocode

Algorithm: Two-Phase PageRank with Travel-Time Self-Loops	
Input:	Graph G , smoothed distribution E_N , parameters μ, β, d
Output:	Predicted traffic $v_{\text{final}} \in \mathbb{R}^N$
1.	Build $N \times N$ transition matrix P using self-loop weights $\text{sw}_i = \mu \cdot \ell_i / s_i$ (Eq. 11)
2.	Build $2N \times 2N$ block matrix M_{2N} (Eq. 18)
3.	Set $E_{2N} \leftarrow [E_N; \mathbf{0}_N]$; initialise $v \leftarrow E_{2N}$
4.	Repeat:
a.	$v_{\text{prev}} \leftarrow v$
b.	$v \leftarrow d \cdot M_{2N}^\top v + (1 - d) \cdot E_{2N}$
Until	$\ v - v_{\text{prev}}\ _1 < 10^{-6}$
5.	$v_{\text{up}} \leftarrow v[1 : N], \quad v_{\text{down}} \leftarrow v[N + 1 : 2N]$
6.	$v_{\text{final}} \leftarrow (v_{\text{up}} + v_{\text{down}}) / \ v_{\text{up}} + v_{\text{down}}\ _1$

Figure 12: Pseudocode for the Two-Phase PageRank power iteration.

6.4 Convergence Analysis

The power iteration converges geometrically at rate $d \cdot \lambda_2(M_{2N})$, where λ_2 is the second-largest eigenvalue of M_{2N} . The convergence bound after k iterations is:

$$\|v^{(k)} - v^*\|_1 \leq (d \cdot \lambda_2)^k \|v^{(0)} - v^*\|_1. \quad (25)$$

In practice, with $d = 0.80$ and our self-loop structure, convergence to 10^{-6} tolerance requires approximately 100–150 iterations for $N = 99,716$ on the SLC network.




Figure 13: Simulated convergence curves (L1 norm of consecutive differences). Higher μ increases the effective spectral gap, leading to faster convergence compared to standard PageRank.

7 Extracting the Final Traffic Distribution

After convergence we have a $2N$ -dimensional stationary vector $v^* = [v_{\text{up}}^*; v_{\text{down}}^*]$. The predicted traffic on physical link i is the total probability mass across both phases:

$$v_{\text{final}}[i] = v_{\text{up}}^*[i] + v_{\text{down}}^*[i], \quad i = 1, \dots, N. \quad (26)$$

We then renormalise to a proper probability distribution:

$$v_{\text{final}} \leftarrow \frac{v_{\text{final}}}{\|v_{\text{final}}\|_1}. \quad (27)$$

This is compared against the normalised ground-truth popularity p_i (from smoothed GPS counts) using the *Mean Relative Error* (MRE):

$$\text{MRE} = \frac{1}{|S|} \sum_{i \in S} \frac{|v_{\text{final}}[i] - p_i|}{p_i}, \quad (28)$$

where S is the evaluation set (either all N links for “Overall MRE”, or the 100 links with highest p_i for “Top-100 MRE”).

8 Parameter Study and Empirical Results

8.1 Baseline: Standard PageRank ($\mu = 0$)

With no self-loops ($\mu = 0$), the algorithm degenerates toward standard PageRank (modulo the two-phase structure). The empirical baseline at ($d = 0.89, \beta = 0.0$):

- Top-100 MRE: 0.7331 — catastrophically high
- Overall MRE: 0.1980

8.2 Effect of the Dwell-Time Parameter μ

We performed a full grid search over the following ranges: $\mu \in \{0, 1, 2, 3, 5, 8, 10, 12, 15, 20\}$; $\beta \in \{0.0, 0.2, 0.5, 0.9\}$; $d \in \{0.80, 0.85, 0.90\}$.

Figure 14 shows the dramatic monotonic improvement in MRE as μ increases, with $d = 0.80$ and $\beta = 0.90$.




Figure 14: Top-100 MRE and Overall MRE as a function of μ for optimal values of d and β . The improvement is monotonic and asymptotes near $\mu = 20$.

Empirical Result

Even a modest $\mu = 1$ reduces Top-100 MRE from 0.6514 to 0.2400—a **63%** improvement—purely by introducing physical time into the model. No other code changes were made.

8.3 Full Parameter Grid Search

Figure 15 visualises the Top-100 MRE across all (μ, β) combinations for $d = 0.80$.




Figure 15: Heatmap of Top-100 MRE for $d = 0.80$ over all tested (μ, β) pairs. The star marks the global optimum at $\mu = 20, \beta = 0.90$.

Table 7: Top-10 best parameter configurations from the full tuning sweep.

μ	d	β	Top-100 MRE	Overall MRE
20	0.80	0.90	0.0317	0.0471
20	0.80	0.50	0.0330	0.0491
20	0.80	0.20	0.0353	0.0526
20	0.80	0.00	0.0393	0.0585
15	0.80	0.90	0.0399	0.0582
15	0.80	0.50	0.0416	0.0607
15	0.80	0.20	0.0445	0.0651
20	0.85	0.90	0.0450	0.0660
20	0.85	0.50	0.0461	0.0676
12	0.80	0.90	0.0475	0.0679

8.4 Sensitivity to Damping Factor d

Figure 16 shows that a lower damping factor $d = 0.80$ consistently outperforms $d = 0.85$ and $d = 0.90$. This is because lower d means more frequent teleportation back to the observed distribution, keeping the simulation grounded in the empirical data.




Figure 16: Sensitivity of Top-100 and Overall MRE to the damping factor d and phase-transition rate β , with $\mu = 20$ fixed.

8.5 Before vs. After: Summary of Improvement




Figure 17: Bar chart comparing baseline (standard PageRank) against the optimised Two-Phase + Self-Loop model. The 95.7% Top-100 improvement is simultaneous with a 76.2% overall improvement.

Empirical Result

Top-100 MRE: 0.7331 $\rightarrow$ **0.0317** (−95.7%)
 Overall MRE: 0.1980 $\rightarrow$ **0.0471** (−76.2%)

Critically, both metrics improve *simultaneously*—confirming that the μ self-loop is a principled physical correction, not an over-fitted patch.

9 Physical Interpretation and Theoretical Connections

9.1 Connection to Continuous-Time Markov Chains

Consider a continuous-time Markov chain (CTMC) where the holding time at state i is exponentially distributed with rate λ_i . In the embedded discrete-time chain, the probability of transitioning to successor j is q_{ij}/λ_i .

Our discrete self-loop construction approximates this with:

$$P_{ii} = \frac{\mu t_i}{\mu t_i + k} \approx 1 - \frac{k}{\mu t_i} \quad \text{for } \mu t_i \gg k. \quad (29)$$

The effective “exit rate” per step is $\approx k/(\mu t_i)$, so the effective holding time (in iterations) is $\approx \mu t_i/k$. This is proportional to the physical traversal time t_i , achieving the desired CTMC approximation.

9.2 Why $\mu = 20$ Is Near-Optimal

For a typical SLC highway ($t_i \approx 36$ s, $k \approx 3$ successors):

$$P_{ii} = \frac{20 \times 36}{20 \times 36 + 3} = \frac{720}{723} \approx 0.9959. \quad (30)$$

For a typical connector ($t_i \approx 2$ s):

$$P_{ii} = \frac{20 \times 2}{20 \times 2 + 3} = \frac{40}{43} \approx 0.930. \quad (31)$$

The ratio of effective holding times $720/40 = 18$ closely matches the physical ratio of traversal times $36/2 = 18$. Beyond $\mu = 20$, self-loops become so strong that convergence slows significantly and returns diminish.

9.3 The Two-Phase Structure as Absorption

The $2N \times 2N$ block matrix with the zero bottom-left block is an *absorbing Markov chain*: once in the Down phase, a walker can never return to the Up phase. The Down phase is therefore a *closed communicating class*.

At the stationary distribution, all mass would eventually be absorbed into the Down phase in the pure transition matrix. The PageRank teleportation (with probability $1 - d$) continuously injects mass back into the Up phase, preventing full absorption and sustaining a nontrivial Up-phase stationary distribution.

9.4 Physical Interpretation of the Two Phases

- **Up phase** captures the “in-transit” traffic: vehicles moving toward their destinations, accumulating on high-throughput arteries according to the self-loop dynamics.

- **Down phase** captures “at-destination” traffic: vehicles in residential streets, parking lots, and local business districts—explained by the neighbourhood-level dispersal that P provides.

The sum $v_{\text{final}} = v_{\text{up}}^* + v_{\text{down}}^*$ aggregates both contributions, giving a complete picture of where vehicles are at any given moment.

10 Complete Mathematical Summary

Table 8 summarises all parameters used in the final model.

Table 8: Complete parameter reference for the Two-Phase PageRank algorithm.

Symbol	Name	Value	Description
N	Network size	99,716	Directed road links
μ	Dwell-time scale	20	Self-loop weight multiplier
β	Phase-drop rate	0.90	Up→Down transition prob.
d	Damping factor	0.80	Follow transition vs. teleport
γ	Smoothing factor	0.26	Graph-diffusion strength
α	Laplace prior	0.10	Pseudo-count for raw counts
ε	Convergence tol.	10^{-6}	Power iteration L1 threshold

The full algorithmic pipeline can be written compactly as:

$$\mathbf{c}_{\text{smooth}} = (1 - \gamma) (\mathbf{c} + \alpha \mathbf{1}) + \gamma \mathbf{P}_{\text{adj}} (\mathbf{c} + \alpha \mathbf{1}) \quad (\text{Smooth})$$

$$E_N = \frac{\mathbf{c}_{\text{smooth}}}{\|\mathbf{c}_{\text{smooth}}\|_1} \quad (\text{Normalise})$$

$$\text{sw}_i = \mu \cdot \ell_i / s_i, \quad P_{ij} = \frac{w_{ij}}{\text{sw}_i + \sum_k w_{ik}} \quad (\text{Build } P)$$

$$M_{2N} = \begin{bmatrix} (1 - \beta)P & \beta I \\ \mathbf{0} & P \end{bmatrix}, \quad E_{2N} = \begin{bmatrix} E_N \\ \mathbf{0} \end{bmatrix} \quad (\text{Build } M_{2N})$$

$$v^* = \lim_{k \rightarrow \infty} (d M_{2N}^\top)^k v^{(0)} + (1 - d) \sum_{k=0}^{\infty} (d M_{2N}^\top)^k E_{2N} \quad (\text{Stationary dist.})$$

$$v_{\text{final}} = \frac{v^*[1 : N] + v^*[N + 1 : 2N]}{\|v^*[1 : N] + v^*[N + 1 : 2N]\|_1} \quad (\text{Final})$$

11 Conclusion

The Two-Phase PageRank with Travel-Time Self-Loops resolves the core deficiency of applying discrete-time random walks to physical road networks. The key contributions are:

1. **Physical dwell time** ($\mu=20$): A single interpretable parameter transforms the transition matrix into a CTMC approximation, allowing realistic traffic concentration to emerge naturally from network topology and link dimensions.

2. **Two-phase structure** ($\beta=0.90$): Separating active driving from local dispersal prevents forward-phase over-diffusion and correctly models where vehicles are parked versus in-transit.
3. **Graph-diffusion pre-processing** ($\gamma=0.26$): Cross-validated spatial smoothing converts the sparse GPS sample into a continuous probability field, eliminating numerical singularities and providing a statistically principled teleportation target.
4. **Simultaneous dual improvement**: The model achieves a **95.7%** reduction in Top-100 MRE and a **76.2%** reduction in Overall MRE *simultaneously*, demonstrating the fix is physically principled and not over-fitted to a single metric.

Empirical Result

By injecting physical reality into the mathematics—traversal time, phase transitions, and spatial smoothing—the algorithm transforms an inaccurate network-flow model into a state-of-the-art traffic prediction system without any machine-learning machinery.